

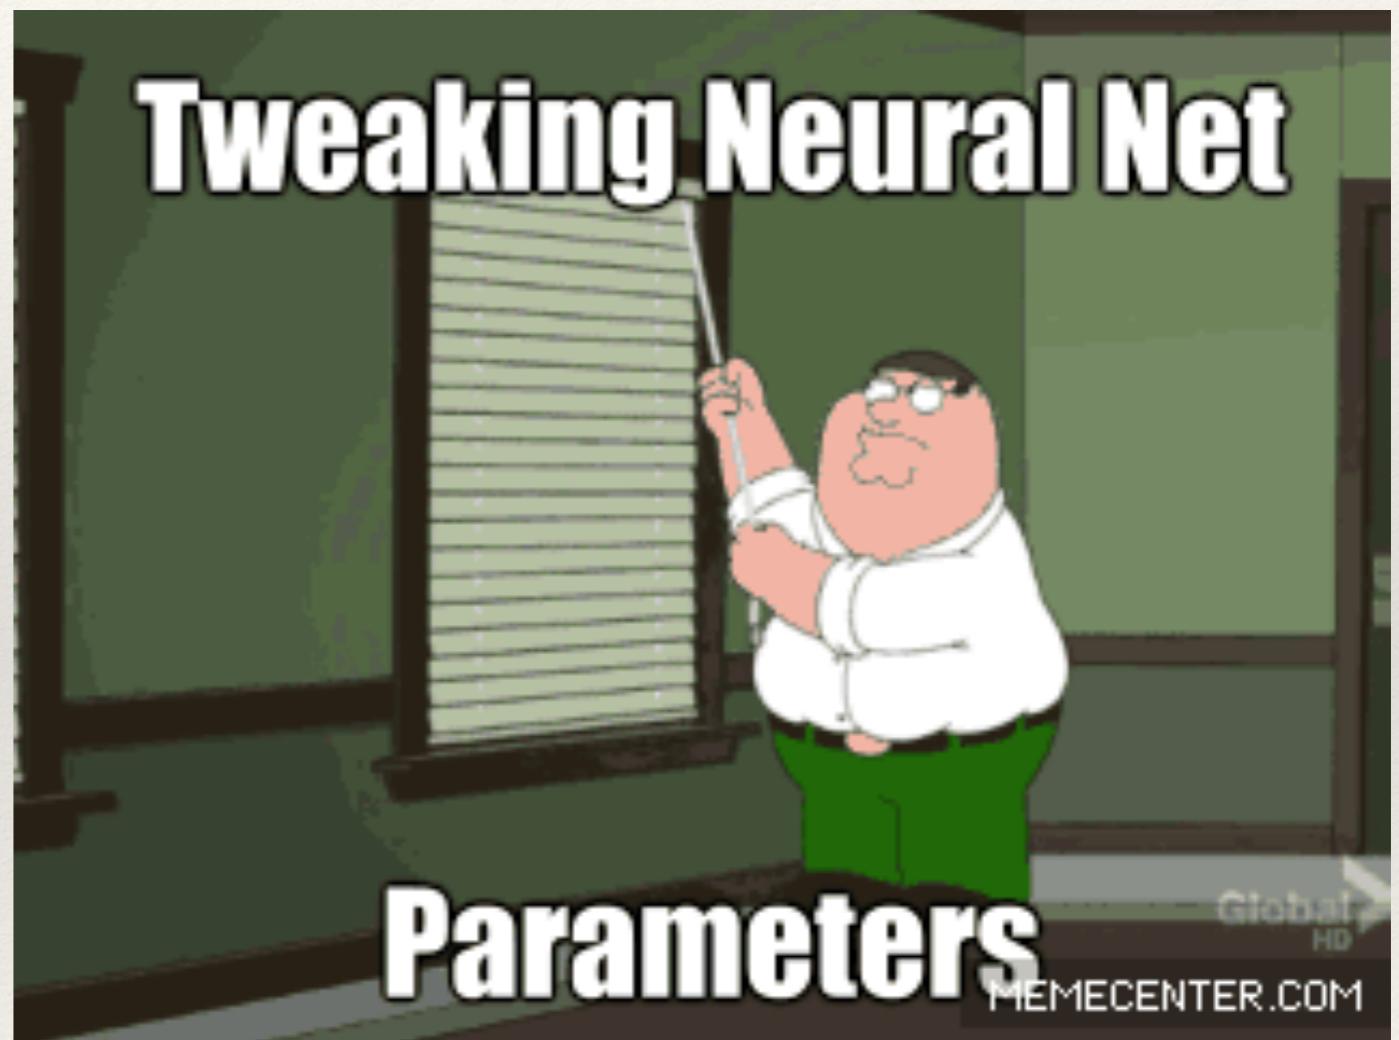
Reza Katebi

Generative Adversarial Networks (GAN)

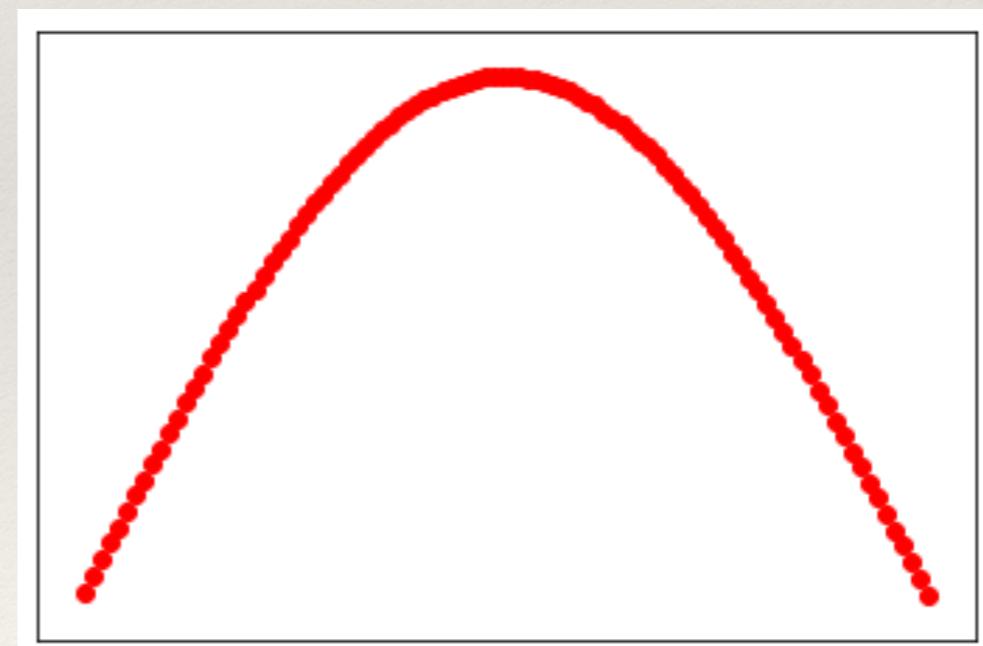
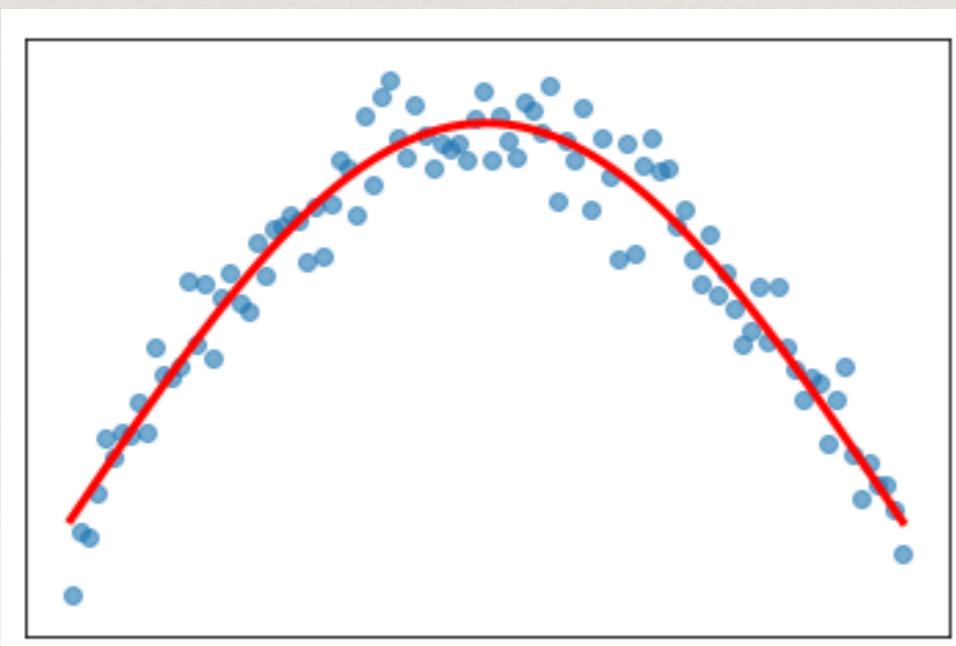
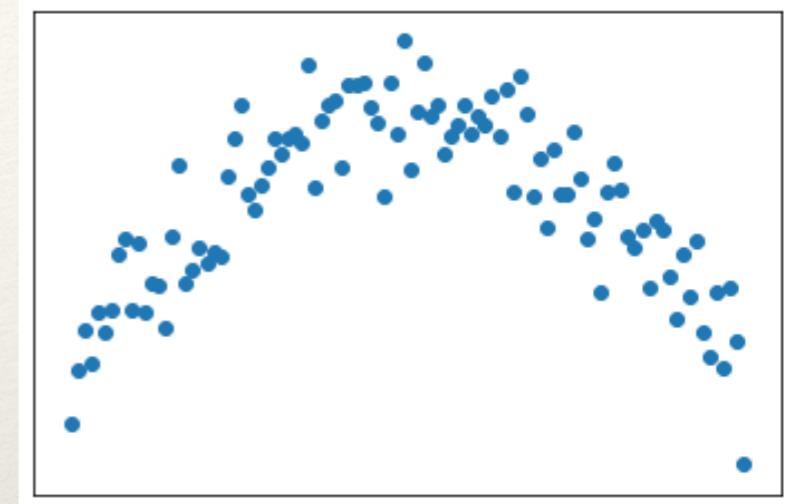
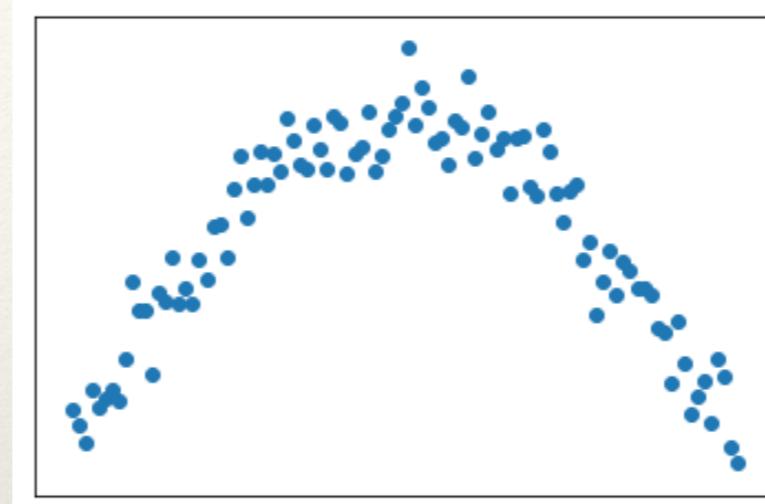
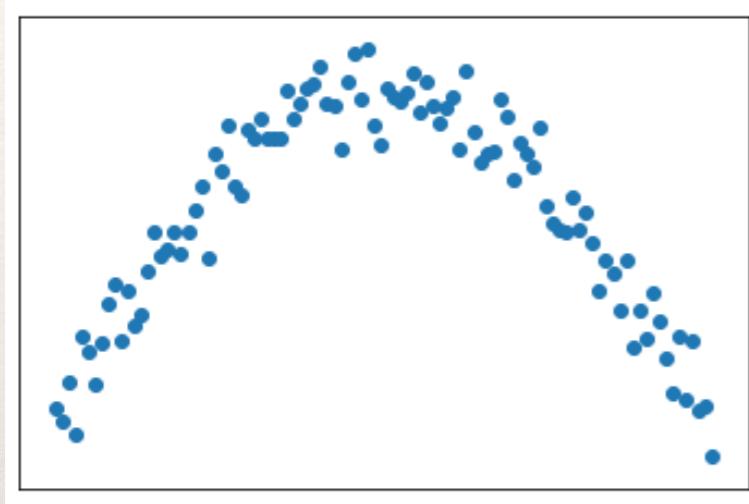
- How Neural Network works and can it be generative?
 - Prisoner Dilemma
 - What is GAN?
 - Implementing it together
-

Neural Networks

- ❖ Trained for classification or regression (Cat and Dog, House Pricing)
- ❖ Somewhere there there is a model of what is cat or dog
- ❖ It only can be used for classification
- ❖ It does not know how to generate new cat since it did not learn how

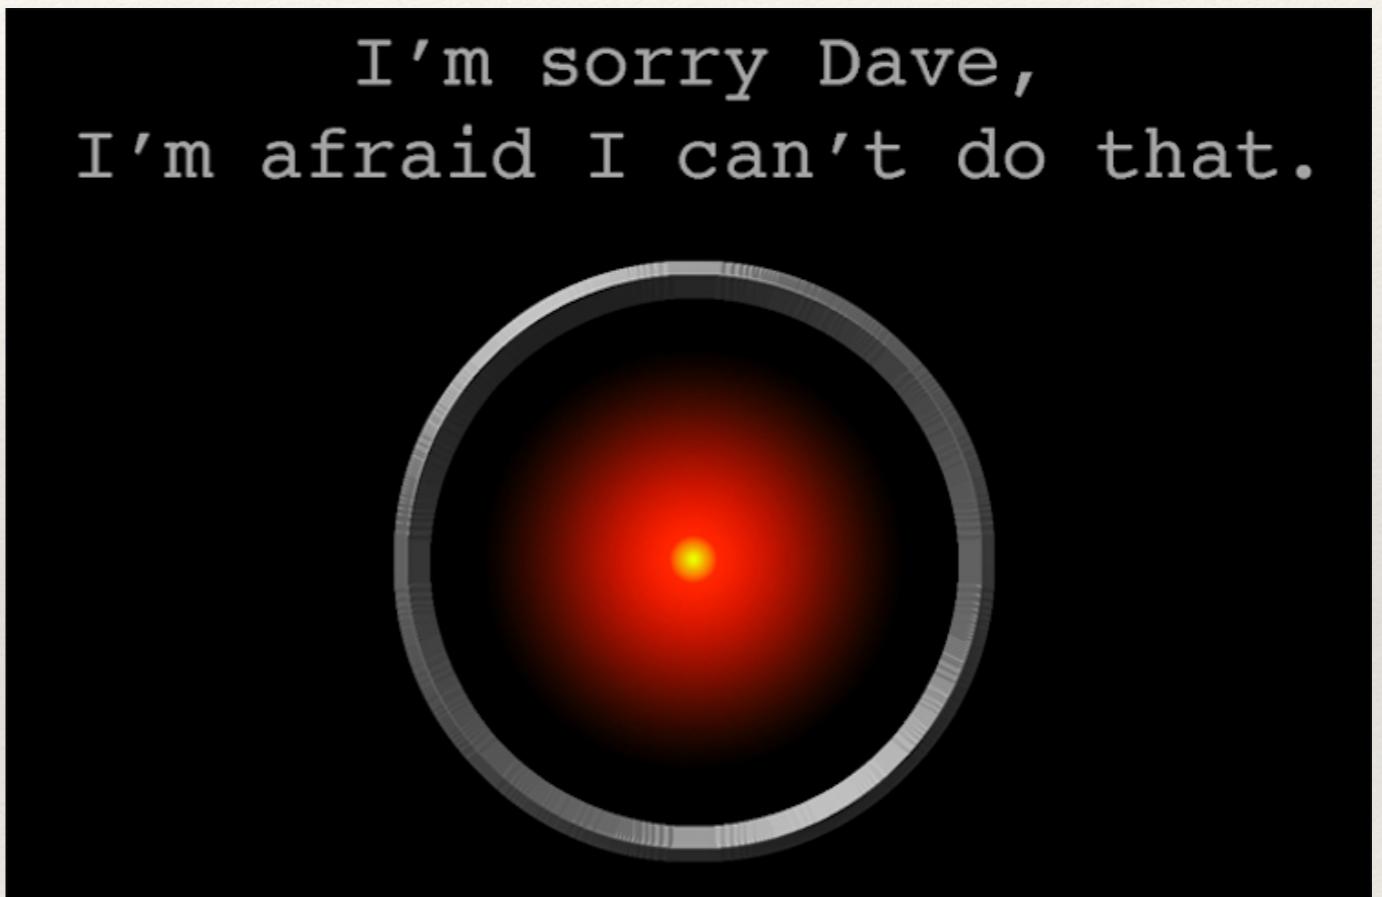


Neural Network



Neural Network

- ❖ Training on MNIST
- ❖ If misclassified 1 and 7
 - ❖ We repeat training for all numbers
- ❖ Humans (children)
 - ❖ If misclassified 1 and 7 and guessed
 - ❖ We focus on 1, 7 only
- ❖ Training where student is failing!
- ❖ You cannot go forever because humans will get tired and quit!
- ❖ You can hammer AI, No feeling yet!

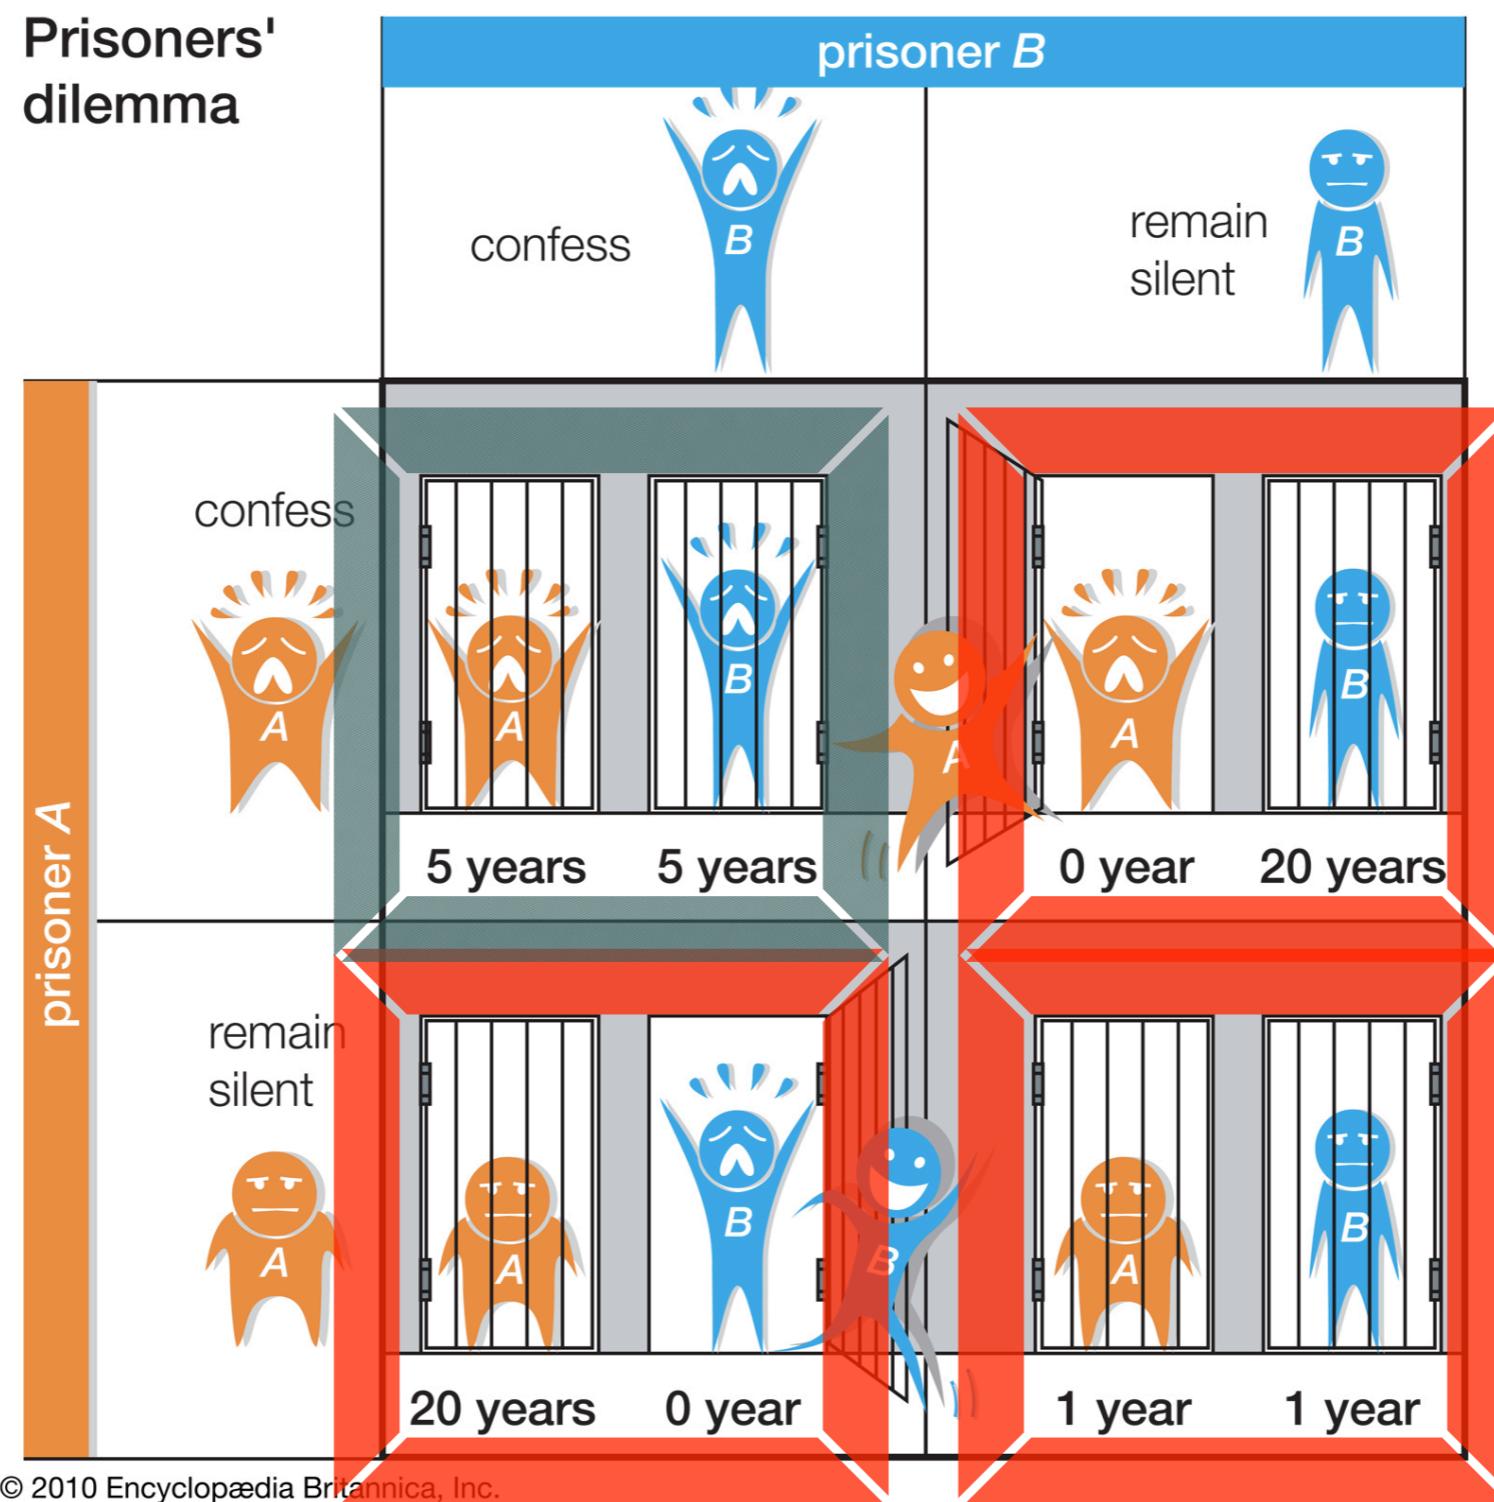


Neural Network

- ❖ Some process that will make the network to output as high as possible error
 - ❖ Spots weakness and forces network to focus on that
 - ❖ Network will resolve that weakness by learning
- ❖ Example: Game playing AI such as Alpha Zero
 - ❖ **19.6 million self played games in FOUR HOURS!!**
- ❖ GAN is something like this.

Zero-Sum Game

Nash Equilibrium



MiniMax Theorem

- ❖ Maximin: Largest value player A can gain not knowing the actions of player B:

$$\bar{V}_A = \max_{a_A} \min_{a_B} v_{a_A}(a_A, a_B)$$

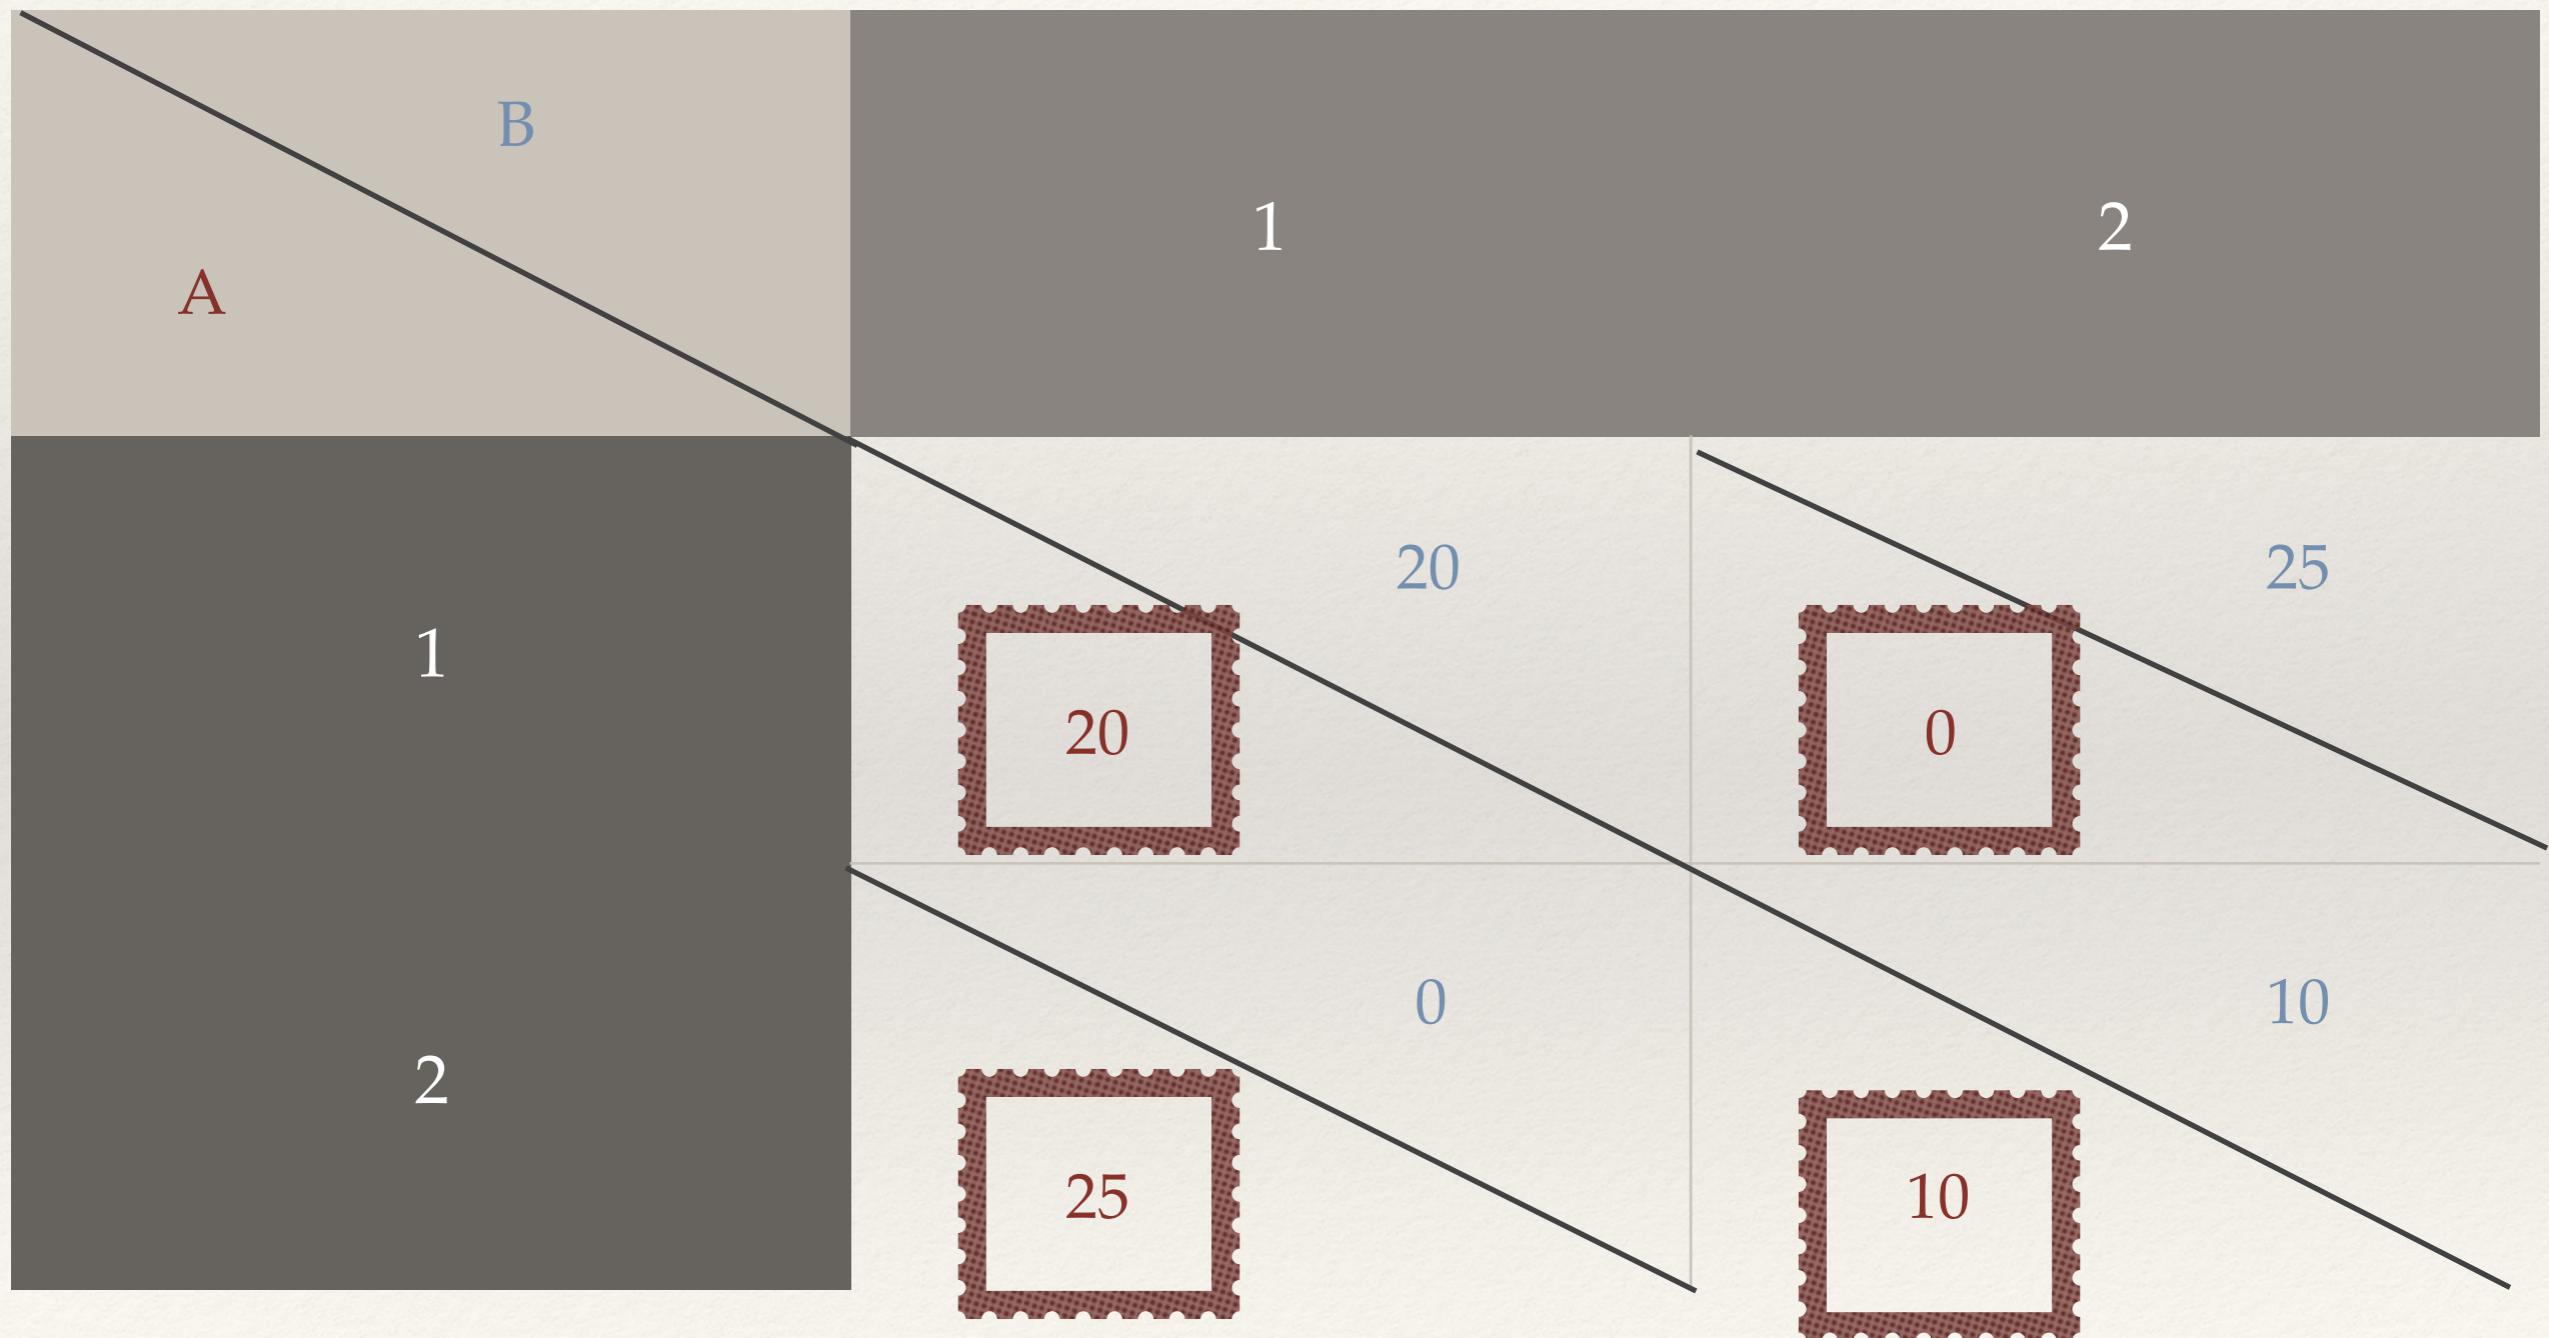
- ❖ Minimax: Smallest value player B can force player A to receive without knowing his actions:

$$\bar{V}_A = \min_{a_B} \max_{a_A} v_{a_A}(a_A, a_B)$$

$\max(0,10) = 10$
A chooses 2

MinMax Example

$\min(25,10) = 10$
B chooses 2



GANs

Algorithm 1 Minibatch stochastic gradient descent training of generative adversarial nets. The number of steps to apply to the discriminator, k , is a hyperparameter. We used $k = 1$, the least expensive option, in our experiments.

for number of training iterations **do**

for k steps **do**

- Sample minibatch of m noise samples $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$ from noise prior $p_g(\mathbf{z})$.
- Sample minibatch of m examples $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$ from data generating distribution $p_{\text{data}}(\mathbf{x})$.
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[\log D(\mathbf{x}^{(i)}) + \log (1 - D(G(\mathbf{z}^{(i)}))) \right].$$

end for

- Sample minibatch of m noise samples $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$ from noise prior $p_g(\mathbf{z})$.
- Update the generator by descending its stochastic gradient:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (1 - D(G(\mathbf{z}^{(i)}))).$$

end for

The gradient-based updates can use any standard gradient-based learning rule. We used momentum in our experiments.

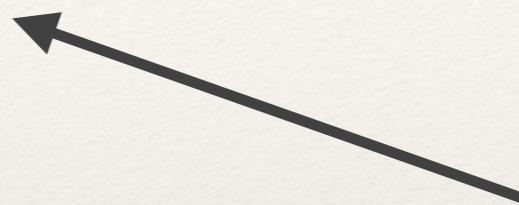
Forger and Investigation Office



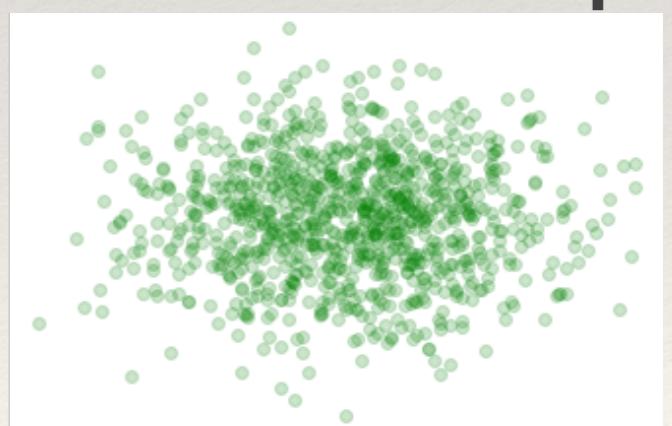
X: Original
Data



D: Investigator
(Discriminator)



G: Forger
(Generator)



Z: Generator
Input

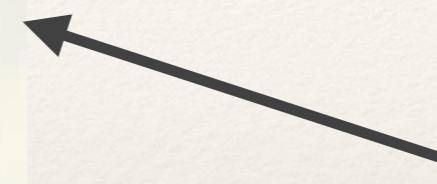
Forger and Investigation Office



X: Original
Data



D: Investigator
(Discriminator)



G: Forger
(Generator)



Z: Generator
Input

Questions?

Why Rectified Linear Unit (ReLU)

- ❖ Computationally less expensive
- ❖ Reduces the likelihood of vanishing gradients
 - ❖ In many layers:
 - ❖ product of many smaller than one gradients gets closer to zero
- ❖ Increases the sparsity

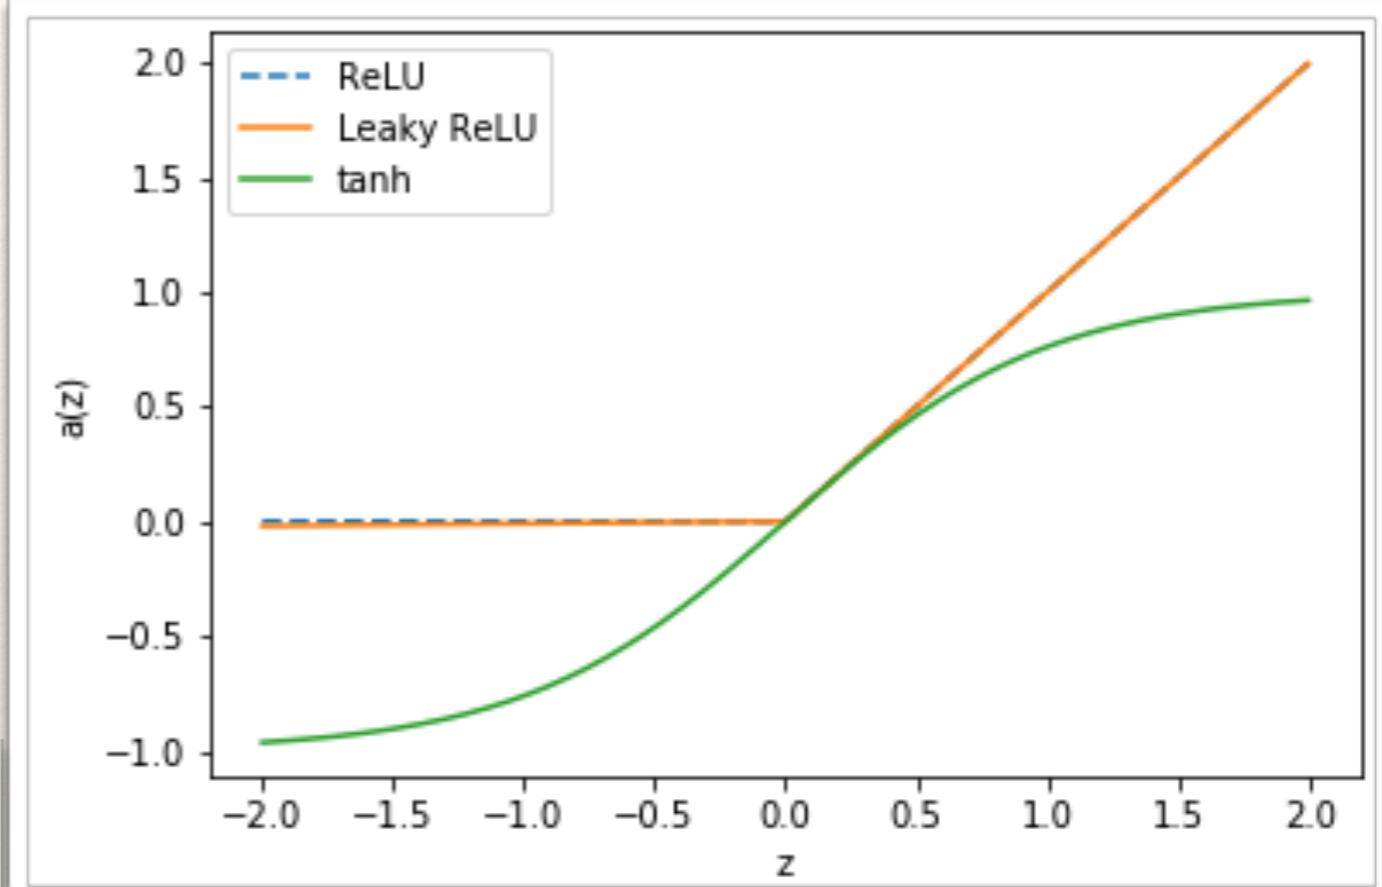
```
def ReLU(x):
    return np.maximum(x, 0)

def leaky_ReLU(x):
    return np.maximum(x, 0.01*x)

def tanh(x):
    return (np.exp(x) - np.exp(-x))/(np.exp(x) + np.exp(-x))

def sigmoid(x):
    return (1. / (1 + np.exp(-x)))

arr = np.random.uniform(-2, 2, 1000)
arr.sort()
plt.plot(arr, ReLU(arr), label = 'ReLU', linestyle = '--')
plt.plot(arr, leaky_ReLU(arr), label = 'Leaky ReLU')
plt.plot(arr, tanh(arr), label = 'tanh')
plt.xlabel('z')
plt.ylabel('a(z)')
plt.legend()
plt.savefig('Activation_Functions.png')
plt.show()
```



Xavier initialization

- ❖ Initialize weights with values that are “not too small” or “not too large”
- ❖ Decreases the likelihood of explosion or vanishing gradients

```
def xavier_init(size):
    in_dim = size[0]
    # xavier variance is 1./in_dim
    # For ReLU it is 2./in_dim
    xavier_stddev = np.sqrt(2./in_dim)
    return Variable(torch.randn(*size) * xavier_stddev)
```

Adaptive Moment Estimation (Adam)

Algorithm 1: Adam, our proposed algorithm for stochastic optimization. See section 2 for details, and for a slightly more efficient (but less clear) order of computation. g_t^2 indicates the elementwise square $g_t \odot g_t$. Good default settings for the tested machine learning problems are $\alpha = 0.001$, $\beta_1 = 0.9$, $\beta_2 = 0.999$ and $\epsilon = 10^{-8}$. All operations on vectors are element-wise. With β_1^t and β_2^t we denote β_1 and β_2 to the power t .

Require: α : Stepsize

Require: $\beta_1, \beta_2 \in [0, 1]$: Exponential decay rates for the moment estimates

Require: $f(\theta)$: Stochastic objective function with parameters θ

Require: θ_0 : Initial parameter vector

$m_0 \leftarrow 0$ (Initialize 1st moment vector)

$v_0 \leftarrow 0$ (Initialize 2nd moment vector)

$t \leftarrow 0$ (Initialize timestep)

while θ_t not converged **do**

$t \leftarrow t + 1$

$g_t \leftarrow \nabla_{\theta} f_t(\theta_{t-1})$ (Get gradients w.r.t. stochastic objective at timestep t)

$m_t \leftarrow \beta_1 \cdot m_{t-1} + (1 - \beta_1) \cdot g_t$ (Update biased first moment estimate)

$v_t \leftarrow \beta_2 \cdot v_{t-1} + (1 - \beta_2) \cdot g_t^2$ (Update biased second raw moment estimate)

$\hat{m}_t \leftarrow m_t / (1 - \beta_1^t)$ (Compute bias-corrected first moment estimate)

$\hat{v}_t \leftarrow v_t / (1 - \beta_2^t)$ (Compute bias-corrected second raw moment estimate)

$\theta_t \leftarrow \theta_{t-1} - \alpha \cdot \hat{m}_t / (\sqrt{\hat{v}_t} + \epsilon)$ (Update parameters)

end while

return θ_t (Resulting parameters)

Kingma, Diederik P., and Jimmy Ba.
"Adam: A method for stochastic optimization."
arXiv preprint arXiv:1412.6980 (2014).

Let's code

Programming in a nutshell

